

# Clustering - Football

## Aim of the Project [🔗](#)

 **App link:** [Streamlit App](#)

- This project aims to cluster football **teams** and **players** based on playing styles and performance metrics.
  - For **teams**: The goal is to group clubs with similar tactical styles, helping analyse their strengths, weaknesses, and unique identities.
  - For **players**: The aim is to identify players with similar profiles, spot standout performers, and understand their strengths and weaknesses in context.

## Data Collection & Cleaning [🔗](#)

### Data Source [🔗](#)

Data was scraped from [FBRef](#) using `pandas.read_html()`.

Example:

```
1 # Example Query to scrape Champions league data table
2 standard_stats = pd.read_html(
3     'https://fbref.com/en/comps/8/stats/Champions-League-Stats',
4     attrs={'id': 'stats_squads_standard_for'}
5 )[0]
```

- The same process was applied to player stats.

### Data Cleaning [🔗](#)

- Scraped data often contains messy formatting, such as:
  - MultiIndex column headers
  - "Unnamed" columns
  - Irregular naming conventions

A generic cleaning script was used across all tables to ensure consistency.

### Flattening MultiIndex Columns [🔗](#)

```
1 standard_stats.columns = [
2     '._'.join(
3         [lvl if not str(lvl).startswith('Unnamed') else '' for lvl in col]
4     ).strip('._')
5     for col in standard_stats.columns.values
6 ]
```

This joins multi-level column headers with underscores and removes "Unnamed" levels.

### Normalizing Column Names [🔗](#)

```
1 def clean_column(col):
2     col = col.strip().lower()
3     col = re.sub(r'\s+', '_', col)
4     col = re.sub(r'_+', '_', col)
5     col = col.strip('_')
```

```

6     return col
7
8     standard_stats.columns = [clean_column(col) for col in standard_stats.columns]

```

This standardizes all column names by:

- Lowercasing
- Replacing spaces with underscores
- Removing duplicate/trailing underscores

### ✓ Cleaning the `squad` Column [🔗](#)

```

1 standard_stats['squad'] = standard_stats['squad'].apply(
2     lambda x: x.split(' ', 1)[1] if isinstance(x, str) and ' ' in x else x
3 )

```

This removes country prefixes (e.g. "ENG Manchester City" → "Manchester City").

## Feature Engineering [🔗](#)

- Most scraped stats are raw totals. To normalize for playing time, I calculated **per 90-minute** metrics, making comparisons fair across teams or players with different minutes played.

$$\text{Per 90 Stat} = \left( \frac{\text{Raw Stat}}{\text{Minutes Played}} \right) \times 90$$

```

1 squad_goalkeeping['performance_sota90'] =
  squad_goalkeeping['performance_sota']/squad_goalkeeping['playing_time_min']*90
2
3 squad_defensive_actions['blocks_blocks90'] =
  squad_defensive_actions['blocks_blocks']/squad_defensive_actions['playing_time_min']
4 squad_defensive_actions['tkl+int90'] =
  squad_defensive_actions['tkl+int']/squad_defensive_actions['playing_time_min']
5 squad_defensive_actions['clr90'] = squad_defensive_actions['clr']/squad_defensive_actions['playing_time_min']
6
7 defensive_features_df = (squad_goalkeeping[['squad', 'performance_ga90', 'performance_sota90']]
8
  .merge(squad_defensive_actions[['squad', 'blocks_blocks90', 'tkl+int90', 'clr90']], on='squad', how='left'))

```

- Additionally, I grouped relevant attributes together based on the feature set I was creating, so example above is defensive features
- Similar features created for players

### 🔒 Defensive Features [🔗](#)

```
Index(['squad', 'performance_ga90', 'performance_sota90', 'blocks_blocks90', 'tkl+int90', 'clr90'])
```

- Goals conceded, Shots conceded, Blocks, Tackles, Clearances

### 🎯 Attacking Features [🔗](#)

```
['squad', 'per_90_minutes_gls', 'per_90_minutes_xg', 'standard_sh/90', 'standard_sot/90']
```

- Goals, xG, Shots, Shots on target

## 🎮 Possession & Passing Features [↗](#)

```
Index(['squad', 'poss', 'touches_att_3rd90', 'total_cmp90', 'total_cmp%', 'prgp90'])
```

Possession, Touches in attacking 3rd, Pass completion, Pass completion %, progressive passes

## Applying K-means Clustering [↗](#)

Now we have our data ready, we will apply K-means clustering.

```
1 # Scale features
2 scaler = StandardScaler()
3 X_scaled = scaler.fit_transform(X)
4
5 # Apply k-means with specified k
6 kmeans = KMeans(n_clusters=k, random_state=42)
7 clusters = kmeans.fit_predict(X_scaled)
```

- First we standardise each feature (mean = 0, std=1) ensures all features contribute equally.
- Then we apply **K-Means** clustering to the scaled data.
- Then assign each row (team/player) to one of `k` clusters.


### 📈 Finding the Optimal K (Elbow Method) [↗](#)

- We can identify the optimal K-Value (Number of Clusters) by using the Elbow Method
- First we need to calculate inertia, this is the sum of squared distances between each data point and its assigned cluster centre

$$\text{Inertia} = \sum_{i=1}^n \|\mathbf{x}_i - \mu_{c_i}\|^2$$

Where:

- $\mathbf{x}_i$  is a data point
- $\mu_{c_i}$  is the centroid of its assigned cluster
- Lower inertia means points are closer to their cluster centers.




- I plotted `inertia` vs `k` and look for the "**elbow point**", this is the optimal K value as any higher K has diminishing returns.
- However I will set up the App so the user can freely choose any cluster amount.

### 🍌 Visualizing Clusters (PCA) [🔗](#)

- Since the data has many dimensions, I used **PCA (Principal Component Analysis)** to reduce it to 2 or 3 dimensions for visualization.

```
1 # PCA for dimensionality reduction to 2D
2 pca = PCA(n_components=2)
3 X_pca = pca.fit_transform(X_scaled)
```

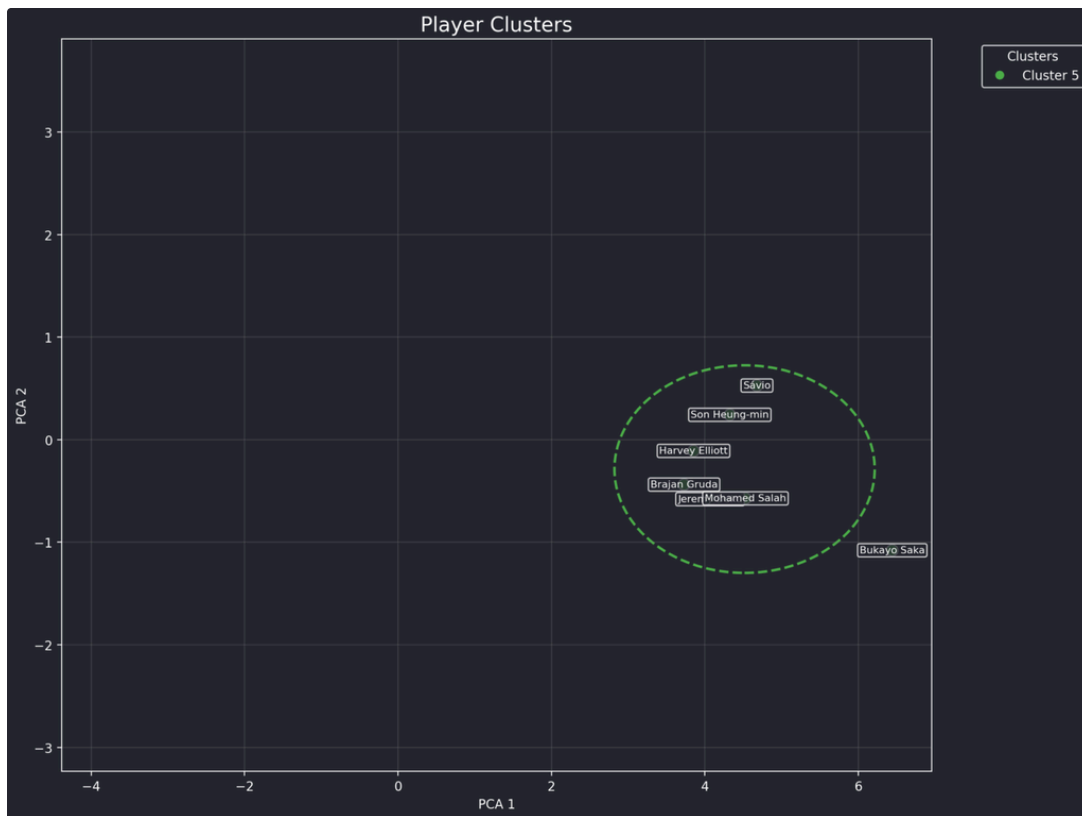
- This preserves structure while simplifying the view.
- The app includes **both 2D and 3D PCA plots** to explore clusters interactively.



- From this PCA plot, we can see the clusters very clearly ( $K=4$ ), and we can even define playstyles based on the cluster averages.
- Additionally in the app, I also implemented a 3D interactive Visualisation for the user to explore. This somewhat resolves the issue of having overlapping clusters in the 2D space.
- I did a similar process for players, any players in the same cluster as each other have similar profiles for that specific feature.

## Example Findings [↗](#)

- All of this analysis has been deployed on a streamlit App, the user can play around with it.
- An interesting finding is Bukayo Saka, when  $K=18$ , features for creativity, we get this output (hide other clusters for clarity)



- The Streamlit app lets users explore clusters dynamically.
- An interesting insight:  
When clustering players by **creativity features** with **K=18**, **Bukayo Saka** appeared at the **edge of his cluster**, showing he's already very unique.
- At **K=19**, Saka was placed in a cluster of his own, suggesting his creative profile is **unmatched** in the dataset.
- Even at K=18, he shares a cluster with elite players like **Salah** and **Son**, indicating elite creativity, but Saka's data shows he's on another level.